

Some slides referred from a lecture note of
George Mason University (Shuochao Yao, Yiwen Xu, Daniel Calzada)

ML Coding Practice
Lecture 02-2

Network Pruning & Quantization

Prof. Jongwon Choi
Chung-Ang University
Fall 2022

Today's Lecture

- Matrix Factorization (SVD)
- Network Pruning
- Network Quantization
- Network Pruning + Quantization

Eigenvalue & Eigenvector

- **Definition**

- $\mathbf{A} \in \mathbb{R}^{n \times n}$ is a square matrix
- Then, $\lambda \in \mathbb{R}$ is an eigenvalue of \mathbf{A} and $\mathbf{x} \in \mathbb{R}^n \setminus \{0\}$ is the corresponding eigenvector of \mathbf{A} if
- $\mathbf{Ax} = \lambda\mathbf{x}$
 - This is called “Eigenvalue equation”

- Conventionally,

- eigenvalues are sorted in descending order

Eigenvalue & Eigenvector

- **Equivalent Statements**

- λ is an eigenvalue of $\mathbf{A} \in \mathbb{R}^{n \times n}$
- There exists an $\mathbf{x} \in \mathbb{R}^n \setminus \{0\}$ with $\mathbf{A}\mathbf{x} = \lambda\mathbf{x}$
- $(\mathbf{A} - \lambda\mathbf{I}_n)\mathbf{x} = 0$ can be solved non-trivially (i.e. $\mathbf{x} \neq 0$)
- $rk(\mathbf{A} - \lambda\mathbf{I}_n) < n$
- $\det(\mathbf{A} - \lambda\mathbf{I}_n) = 0$

Properties of e-val. & e-vec.

- **Properties**

- \mathbf{A} and \mathbf{A}^T have the same e-val, but not necessarily the same e-vec
- The eigenspace E_λ is the null space of $\mathbf{A} - \lambda \mathbf{I}_n$
 - pf. $\mathbf{Ax} = \lambda \mathbf{x} \Rightarrow \mathbf{Ax} - \lambda \mathbf{x} = 0 \Rightarrow (\mathbf{A} - \lambda \mathbf{I})\mathbf{x} = 0 \Rightarrow \mathbf{x} \in \ker(\mathbf{A} - \lambda \mathbf{I})$
- Similar matrices have the same e-val
 - Thus, e-val is independent of the choice of basis of the transformation matrices
 - e-val, determinant, trace are the key characteristic parameters of a linear mapping, which are invariant under basis changes!
- Symmetric, positive definite matrices always have positive, real e-vals.

Linear independence & e-val

- The eigenvectors x_1, \dots, x_n of a matrix $A \in \mathbb{R}^{n \times n}$ with n distinct eigenvalues $\lambda_1, \dots, \lambda_n$ are linearly independent
- **Defective**
 - A square matrix $A \in \mathbb{R}^{n \times n}$ is defective if it possesses fewer than n linearly independent e-vects.
 - Remark!
 - A non-defective matrix does not necessarily require n distinct e-vals.

Determinant & Trace & e-vals

- The determinant of a matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ is the product of its e-vals:
 - $\det(\mathbf{A}) = \prod_{i=1}^n \lambda_i$
 - where $\lambda \in \mathbb{C}$ are (possibly repeated) e-vals of \mathbf{A}
- The trace of $\mathbf{A} \in \mathbb{R}^{n \times n}$ is the sum of its e-vals:
 - $\text{tr}(\mathbf{A}) = \sum_{i=1}^n \lambda_i$
 - where $\lambda \in \mathbb{C}$ are (possibly repeated) e-vals of \mathbf{A}

Cholesky Decomposition

- **Necessity**

- The square-root operation in matrix
- In the case of single number, the square-root operation is only for the positive number

- **Definition**

- A square-root equivalent operation
- Only possible for symmetric, positive definite matrices

Cholesky Decomposition

- **Definition**

- A square-root equivalent operation
- Only possible for symmetric, positive definite matrices
- For a symmetric, positive definite matrix \mathbf{A} ,
- \mathbf{A} can be factorized into a product $\mathbf{A} = \mathbf{L}\mathbf{L}^T$
- , where \mathbf{L} is a lower-triangular matrix with positive diagonal elements

- $$\begin{bmatrix} a_{11} & \dots & a_{1n} \\ \dots & \dots & \dots \\ a_{n1} & \dots & a_{nn} \end{bmatrix} = \begin{bmatrix} l_{11} & \dots & 0 \\ \dots & \dots & \dots \\ l_{n1} & \dots & l_{nn} \end{bmatrix} \begin{bmatrix} l_{11} & \dots & l_{n1} \\ \dots & \dots & \dots \\ 0 & \dots & l_{nn} \end{bmatrix}$$

Cholesky Decomposition

- **Interesting Property**

- We can compute determinants very efficiently!
- $\det(\mathbf{A}) = \det(\mathbf{L}) \det(\mathbf{L}^T) = \det(\mathbf{L})^2$
- Here, since \mathbf{L} is a triangular matrix,
- $\det(\mathbf{L}) = \prod_{i=1}^n l_{ii}$
- Thus, $\det(\mathbf{A}) = \det(\mathbf{L})^2 = \prod_{i=1}^n l_{ii}^2$

Diagonal Matrix

- **Definition**

- A matrix has value 0 on all off-diagonal elements

- $$D = \begin{bmatrix} c_1 & \dots & 0 \\ \dots & \dots & \dots \\ 0 & \dots & c_n \end{bmatrix}$$

- **Notation**

- $$D = \begin{bmatrix} c_1 & \dots & 0 \\ \dots & \dots & \dots \\ 0 & \dots & c_n \end{bmatrix} = \text{diag}([c_1, \dots, c_n])$$

Diagonal Matrix

- **Properties**

- Fast computation of determinants, powers, and inverses!

- Determinant

- $\det(\mathbf{D}) = \prod_{i=1}^n c_i$

- Power

- $\mathbf{D}^k = \text{diag}([c_1^k, \dots, c_n^k])$

- Inverse (Only if only the non-zero values are diagonal)

- $\mathbf{D}^{-1} = \text{diag}\left(\left[\frac{1}{c_1}, \dots, \frac{1}{c_n}\right]\right)$

Diagonalizable

- **e-val & e-vec & Diagonalizable**

- $\mathbf{A} \in \mathbb{R}^{n \times n}$

- $\lambda_1, \dots, \lambda_n \in \mathbb{R}^1$

- $\mathbf{p}_1, \dots, \mathbf{p}_n \in \mathbb{R}^n$

- Then, $\mathbf{AP} = \mathbf{PD}$

- iff λ_i are e-val of \mathbf{A} and \mathbf{p}_i are the corresponding e-vec

Eigendecomposition

- A square matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ can be factored into
 - $\mathbf{A} = \mathbf{P}\mathbf{D}\mathbf{P}^{-1}$
 - , where $\mathbf{P} \in \mathbb{R}^{n \times n}$, $\mathbf{D} = \text{diag}([\lambda_1, \dots, \lambda_n])$, and λ_i are the e-val of \mathbf{A}
- iff the e-val of \mathbf{A} form a basis of \mathbb{R}^n

Computation Efficiency

- **Power of A**

- $\mathbf{A}^k = (\mathbf{P}\mathbf{D}\mathbf{P}^{-1})^k = \mathbf{P}\mathbf{D}^k\mathbf{P}^{-1}$

- **Determinant of A**

- $\det(\mathbf{A}) = \det(\mathbf{P}\mathbf{D}\mathbf{P}^{-1}) = \det(\mathbf{P}) \det(\mathbf{D}) \det(\mathbf{P}^{-1}) = \det(\mathbf{D}) = \prod_{i=1}^n d_{ii}$

Limitation of e-val decomposition

Only works for the square matrix!

- **Cholesky Decomposition**

- For a symmetric, positive definite matrix A ,
- A can be factorized into a product $A = LL^T$
- , where L is a lower-triangular matrix with positive diagonal elements

- $$\begin{bmatrix} a_{11} & \dots & a_{1n} \\ \dots & \dots & \dots \\ a_{n1} & \dots & a_{nn} \end{bmatrix} = \begin{bmatrix} l_{11} & \dots & 0 \\ \dots & \dots & \dots \\ l_{n1} & \dots & l_{nn} \end{bmatrix} \begin{bmatrix} l_{11} & \dots & l_{n1} \\ \dots & \dots & \dots \\ 0 & \dots & l_{nn} \end{bmatrix}$$

- **Diagonalization**

- A square matrix $A \in \mathbb{R}^{n \times n}$ can be factored into
 - $A = PDP^{-1}$
 - , where $P \in \mathbb{R}^{n \times n}$, $D = \text{diag}([\lambda_1, \dots, \lambda_n])$, and λ_i are the e-val of A
- iff the e-val of A form a basis of \mathbb{R}^n

SVD Theorem

- $\mathbf{A}_{m \times n}$ is a rectangular matrix of rank $r \in [0, \min(m, n)]$
- The SVD of \mathbf{A} is a decomposition of the form
 - $\mathbf{A}_{m \times n} = \mathbf{U}_{m \times m} \mathbf{\Sigma}_{m \times n} \mathbf{V}_{n \times n}^T$
 - $\mathbf{U}_{m \times m}$: orthogonal matrix with column vectors \mathbf{u}_i
 - $\mathbf{V}_{n \times n}$: orthogonal matrix with column vectors \mathbf{v}_i
 - $\mathbf{\Sigma}_{m \times n}$: $\Sigma_{ii} = \sigma_i \geq 0$ and $\Sigma_{ij} = 0, i \neq j$

SVD Theorem

- $\mathbf{A}_{m \times n} = \mathbf{U}_{m \times m} \mathbf{\Sigma}_{m \times n} \mathbf{V}_{n \times n}^T$
- $\mathbf{U}_{m \times m}$: orthogonal matrix with column vectors \mathbf{u}_i
 - \mathbf{u}_i are called the left-singular vectors
- $\mathbf{V}_{n \times n}$: orthogonal matrix with column vectors \mathbf{v}_i
 - \mathbf{v}_i are called the right-singular vectors
- $\mathbf{\Sigma}_{m \times n}$: $\Sigma_{ii} = \sigma_i \geq 0$ and $\Sigma_{ij} = 0, i \neq j$
 - σ_i are called the “singular values”
 - Conventionally, the singular values are sorted, i.e. $\sigma_1 \geq \sigma_2 \geq \dots \geq 0$
 - Unique!!

SVD – Mathematical Meaning

- SVD vs. Eigendecomposition
 - Eigendecomposition: $\mathbf{S} = \mathbf{S}^T = \mathbf{P}\mathbf{D}\mathbf{P}^T$ (Of course, \mathbf{S} is SPD)
 - SVD: $\mathbf{S} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$
- Then,
 - When $\mathbf{U} = \mathbf{P} = \mathbf{V}$, $\mathbf{D} = \mathbf{\Sigma}$
 - SVD of SPD matrices is their eigendecomposition!

SVD – Mathematical Meaning

- What's the direct connection between eigendecomposition and SVD
 - Without any constraint (i.e. SPD)
- $\mathbf{A}^T \mathbf{A} \in \mathbb{R}^{n \times n}$ is symmetric, positive semi-definite matrix
 - Any matrix $A \in \mathbb{R}^{m \times n}$
- The symmetric, positive semi-definite matrix can be diagonalized as:
 - $\mathbf{A}^T \mathbf{A} = \mathbf{P} \mathbf{D} \mathbf{P}^T = \mathbf{P} \begin{bmatrix} \lambda_1 & \dots & 0 \\ \dots & \dots & \dots \\ 0 & \dots & \lambda_n \end{bmatrix} \mathbf{P}^T$
 - where \mathbf{P} is an orthogonal matrix composed of the orthonormal eigenbasis
 - $\lambda_i \geq 0$

SVD – Mathematical Meaning

- $\mathbf{A}^T \mathbf{A} \in \mathbb{R}^{n \times n}$ is symmetric, positive semi-definite matrix
- $\mathbf{A}^T \mathbf{A} = \mathbf{P} \mathbf{D} \mathbf{P}^T = \mathbf{P} \begin{bmatrix} \lambda_1 & \dots & 0 \\ \dots & \dots & \dots \\ 0 & \dots & \lambda_n \end{bmatrix} \mathbf{P}^T$
- $\mathbf{A}^T \mathbf{A} = (\mathbf{U} \mathbf{\Sigma} \mathbf{V}^T)^T (\mathbf{U} \mathbf{\Sigma} \mathbf{V}^T) = \mathbf{V} \mathbf{\Sigma}^T \mathbf{U}^T \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T$
- Here, remark that \mathbf{U} , \mathbf{V} are the orthonormal matrices
- Thus, $\mathbf{U}^T \mathbf{U} = \mathbf{I}$

SVD – Mathematical Meaning

- $\mathbf{A}^T \mathbf{A} \in \mathbb{R}^{n \times n}$ is symmetric, positive semi-definite matrix
- $\mathbf{A}^T \mathbf{A} = \mathbf{P} \mathbf{D} \mathbf{P}^T = \mathbf{P} \begin{bmatrix} \lambda_1 & \dots & 0 \\ \dots & \dots & \dots \\ 0 & \dots & \lambda_n \end{bmatrix} \mathbf{P}^T$
- $\mathbf{A}^T \mathbf{A} = \mathbf{V} \mathbf{\Sigma}^T \mathbf{U}^T \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T = \mathbf{V} \mathbf{\Sigma}^T \mathbf{\Sigma} \mathbf{V}^T = \mathbf{V} \begin{bmatrix} \sigma_1^2 & \dots & 0 \\ \dots & \dots & \dots \\ 0 & \dots & \sigma_n^2 \end{bmatrix} \mathbf{V}^T$
- $\mathbf{P}^T = \mathbf{V}^T, \sigma_i^2 = \lambda_i$
- Thus, eigenvectors of $\mathbf{A}^T \mathbf{A}$ are the right-singular vectors \mathbf{V} of \mathbf{A}
- eigenvalues of $\mathbf{A}^T \mathbf{A}$ are the squared singular values of $\mathbf{\Sigma}$

SVD – Mathematical Meaning

- $\mathbf{AA}^T \in \mathbb{R}^{m \times m}$ is symmetric, positive semi-definite matrix
- $\mathbf{AA}^T = \mathbf{S}\mathbf{D}\mathbf{S}^T = \mathbf{S} \begin{bmatrix} \lambda_1 & \dots & 0 \\ \dots & \dots & \dots \\ 0 & \dots & \lambda_n \end{bmatrix} \mathbf{S}^T$
- $\mathbf{AA}^T = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T\mathbf{V}\mathbf{\Sigma}^T\mathbf{U}^T = \mathbf{U}\mathbf{\Sigma}\mathbf{\Sigma}^T\mathbf{U}^T = \mathbf{U} \begin{bmatrix} \sigma_m^2 & \dots & 0 \\ \dots & \dots & \dots \\ 0 & \dots & \sigma_m^2 \end{bmatrix} \mathbf{U}^T$
- $\mathbf{S}^T = \mathbf{U}^T, \sigma_i^2 = \lambda_i$
- Thus, eigenvectors of \mathbf{AA}^T are the left-singular vectors \mathbf{U} of \mathbf{A}
- Eigenvalues of \mathbf{AA}^T are the squared singular values of $\mathbf{\Sigma}$

SVD – Mathematical Meaning

- $\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$
- Eigenvectors of $\mathbf{A}\mathbf{A}^T$ are the left-singular vectors \mathbf{U} of \mathbf{A}
- Eigenvectors of $\mathbf{A}^T\mathbf{A}$ are the right-singular vectors \mathbf{V} of \mathbf{A}
- Eigenvalues of $\mathbf{A}\mathbf{A}^T$ and $\mathbf{A}^T\mathbf{A}$ are the squared singular values of $\mathbf{\Sigma}$

SVD – Matrix Approximation

- $\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T = \sum_{i=1}^r \sigma_i \mathbf{u}_i \mathbf{v}_i^T \equiv \sum_{i=1}^r \sigma_i \mathbf{A}_i$
- Remark that $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_r \geq 0$
- Thus, we can easily find the rank-k approximation by:
 - $\hat{\mathbf{A}}(k) = \sum_{i=1}^k \sigma_i \mathbf{u}_i \mathbf{v}_i^T = \sum_{i=1}^k \sigma_i \mathbf{A}_i$
- This is called “Low-rank Approximation”

SVD & Pseudo-inverse Matrix

- $A = U\Sigma V^T$
- Basically, they satisfies
 - U : Orghogonal
 - $AA^T = U(\Sigma\Sigma^T)U^T$
 - $U^{-1} = U^T$
 - V : Orghogonal
 - $A^TA = V(\Sigma^T\Sigma)V^T$
 - $V^{-1} = V^T$
- $A^+ = (A^TA)^{-1}A^T = (V(\Sigma^T\Sigma)V^T)^{-1}V\Sigma^TU^T = V(\Sigma^{-1}\Sigma^{-T})V^TV\Sigma^TU^T$
- Thus, $A^+ = V\Sigma^{-1}U^T$

Fully Connected Layers: Singular Value Decomposition

- Most weights are in the fully connected layers (according to Denton et al.)
- $W = USV^T$
 - $W \in \mathbb{R}^{m \times k}, U \in \mathbb{R}^{m \times m}, S \in \mathbb{R}^{m \times k}, V^T \in \mathbb{R}^{k \times k}$
- S is diagonal, decreasing magnitudes along the diagonal

$$\begin{pmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & A & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{pmatrix} = \begin{pmatrix} \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & U & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \end{pmatrix} \begin{pmatrix} w_1 & & & & \\ & w_2 & & & \\ & & w_3 & & \\ & & & & \\ & & & & \end{pmatrix} \begin{pmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & V^T & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{pmatrix}$$

Singular Value Decomposition

- By only keeping the t singular values with largest magnitude:

- $\tilde{W} = \tilde{U}\tilde{S}\tilde{V}^\top$

- $\tilde{W} \in \mathbb{R}^{m \times k}, \tilde{U} \in \mathbb{R}^{m \times t}, \tilde{S} \in \mathbb{R}^{t \times t}, \tilde{V}^\top \in \mathbb{R}^{t \times k}$

- $\text{Rank}(\tilde{W}) = t$

$$\begin{pmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & A & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{pmatrix} = \begin{pmatrix} \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & U & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \end{pmatrix} \begin{pmatrix} \mathcal{W}_1 & & & & \\ & \mathcal{W}_2 & & & \\ & & \mathcal{W}_3 & & \\ & & & & \\ & & & & \end{pmatrix} \begin{pmatrix} \cdot & \cdot & \cdot \\ \cdot & V^T & \cdot \\ \cdot & \cdot & \cdot \end{pmatrix}$$

SVD: Compression

- $W = USV^\top, W \in \mathbb{R}^{m \times k}, U \in \mathbb{R}^{m \times m}, S \in \mathbb{R}^{m \times k}, V^\top \in \mathbb{R}^{k \times k}$
- $\tilde{W} = \tilde{U}\tilde{S}\tilde{V}^\top, \tilde{W} \in \mathbb{R}^{m \times k}, \tilde{U} \in \mathbb{R}^{m \times t}, \tilde{S} \in \mathbb{R}^{t \times t}, \tilde{V}^\top \in \mathbb{R}^{t \times k}$
- Storage for W : $O(mk)$
- Storage for \tilde{W} : $O(mt + t + tk)$
- Compression Rate: $O\left(\frac{mk}{t(m+k+1)}\right)$
- Theoretical error: $\|A\tilde{W} - AW\|_F \leq s_{t+1}\|A\|_F$

Gong, Yunchao, et al. "Compressing deep convolutional networks using vector quantization." *arXiv preprint arXiv:1412.6115* (2014).

SVD: Compression Results

- Trained on ImageNet 2012 database, then compressed
- 5 convolutional layers, 3 fully connected layers, softmax output layer

Approximation method	Number of parameters	Approximation hyperparameters	Reduction in weights	Increase in error
Standard FC	$N M$			
FC layer 1: Matrix SVD	$N K + K M$	$K = 250$ $K = 950$	$13.4 \times$ $3.5 \times$	0.8394% 0.09%
FC layer 2: Matrix SVD	$N K + K M$	$K = 350$ $K = 650$	$5.8 \times$ $3.14 \times$	0.19% 0.06%
FC layer 3: Matrix SVD	$N K + K M$	$K = 250$ $K = 850$	$8.1 \times$ $2.4 \times$	0.67% 0.02%

K refers to rank of approximation, t in the previous slides.

Denton, Emily L., et al. "Exploiting linear structure within convolutional networks for efficient evaluation." *Advances in Neural Information Processing Systems*. 2014.


SVD: Side Benefits

- Reduced memory footprint
 - Reduced in the dense layers by 5-13x
- Speedup: $A\tilde{W}$, $A \in \mathbb{R}^{n \times m}$, computed in $O(nmt + nt^2 + ntk)$ instead of $O(nmk)$
 - Speedup factor is $O\left(\frac{mk}{t(m+t+k)}\right)$
- Regularization
 - “Low-rank projections effectively decrease number of learnable parameters, suggesting that they might improve generalization ability.”
 - Paper applies SVD after training

Denton, Emily L., et al. "Exploiting linear structure within convolutional networks for efficient evaluation." *Advances in Neural Information Processing Systems*. 2014.

Convolutions: Matrix Multiplication

Most time is spent in the convolutional layers




$$F(x, y) = I * W$$


<http://stackoverflow.com/questions/15356153/how-do-convolution-matrices-work>

Flattened Convolutions

- Replace $c \times y \times x$ convolutions with $c \times 1 \times 1$, $1 \times y \times 1$, and $1 \times 1 \times x$ convolutions



(a) 3D convolution



(b) 1D convolutions over different directions

Flattened Convolutions

$$\begin{aligned}\hat{F}(x, y) &= I * \hat{W} \\ &= \sum_{x'=1}^X \left(\sum_{y'=1}^Y \left(\sum_{c=1}^C I(c, x - x', y - y') \alpha(c) \right) \beta(y') \right) \gamma(x')\end{aligned}$$

$\alpha \in \mathbb{R}^C, \beta \in \mathbb{R}^Y, \gamma \in \mathbb{R}^X$

- Compression and Speedup:
 - Parameter reduction: $O(XYC)$ to $O(X + Y + C)$
 - Operation reduction: $O(mnCXY)$ to $O(mn(C + X + Y))$ (where $W_f \in \mathbb{R}^{m \times n}$)

Jin, Jonghoon, Aysegul Dundar, and Eugenio Culurciello. "Flattened convolutional neural networks for feedforward acceleration." *arXiv preprint arXiv:1412.5474* (2014).

Flattening = MF



$$\begin{aligned}\hat{F}(x, y) &= \sum_{\substack{x=1 \\ \bar{X}}}^X \sum_{\substack{y'=1 \\ \bar{Y}}}^Y \sum_{\substack{c=1 \\ \bar{C}}}^C I(c, x - x', y - y') \alpha(c) \beta(y') \gamma(x') \\ &= \sum_{x=1}^X \sum_{y'=1}^Y \sum_{c=1}^C I(c, x - x', y - y') \hat{W}(c, x', y')\end{aligned}$$

- $\hat{W} = \alpha \otimes \beta \otimes \gamma, \text{Rank}(\hat{W}) = 1$
- $\hat{W}_S = \sum_{k=1}^K \alpha_k \otimes \beta_k \otimes \gamma_k, \text{Rank } K$
- SVD: Can reconstruct the original matrix as $A = \sum_{k=1}^K w_k u_k \otimes v_k$


Denton, Emily L., et al. "Exploiting linear structure within convolutional networks for efficient evaluation." *Advances in Neural Information Processing Systems*. 2014.

Flattening: Speedup Results

- 3 convolutional layers (5x5 filters) with 96, 128, and 256 channels
- Used stacks of 2 rank-1 convolutions



(c) Backpropagation on CPU



(a) Feedforward on CPU

Jin, Jonghoon, Aysegul Dundar, and Eugenio Culurciello. "Flattened convolutional neural networks for feedforward acceleration." *arXiv preprint arXiv:1412.5474* (2014).

Today's Lecture

- Matrix Factorization (SVD)
- Network Pruning
- Network Quantization
- Network Pruning + Quantization